

CSI3131 – Operating Systems

Tutorial 3 – Threads

1. What are two differences between user-level threads and kernel-level threads? Under what circumstances is one type better than the other?
2. What resources are used when a thread is created? How do they differ from those used when a process is created?
3. Describe the relationship between user and kernel threads for each of the following models: Many-to-One, One-to-one, and Many-to-Many. In particular, indicate where scheduling is done in each of these models.
4. The program shown below uses Pthreads. What would be output from the program at LINE C and at LINE P? Explain your answer.

```
#include <pthread.h>
#include <stdio.h>

int value = 0;
void *runner(void *param); /* function for the thread */

int main(int argc, char *argv[])
{
    int pid;
    pthread_t tid;
    pthread_attr_t attr;

    pid = fork();

    if(pid == 0) /* the child */
    {
        pthread_attr_init(&attr);
        pthread_create(&tid, &attr, runner, NULL);
        pthread_join(tid, NULL);
        printf("CHILD: value = %d", value); /* LINE C */
    }
    else if(pid > 0) /* the parent */
    {
        wait(NULL);
        printf("PARENT: value = %d\n"); /* LINE P */
    }
}

void *runner(void *param)
{
    value = 5;
    pthread_exit(0);
}
```

5. Consider that two processes start at time $t=0$. Process A operates with a single thread, while Process B runs a multi-threading program that initiates three threads. Thread T1 starts at time $t=0$, thread T2 starts after T1 has run (i.e. assigned to the CPU) for 19 time units (t.u.), and thread T3 starts after T1 has run for 27 t.u. (that is, T1 invokes the T2 and T3 threads). Process A and each of the threads of Process B run a series of CPU bursts separated by I/O operations requested from the OS. Assume that each I/O operation takes 52 time units to complete. The execution of Process A and each thread of Process B occur as follows:

Process A (no multithreading used)

55 t.u. CPU burst, I/O request, 50 t.u. CPU burst, I/O request, 30 t.u. CPU burst, I/O request

Process B (multi-threaded)


T1: 300 t.u. CPU burst (CPU bound thread)

T2: 12 t.u. CPU burst, followed by I/O request (repeated 5 time) (I/O bound thread)

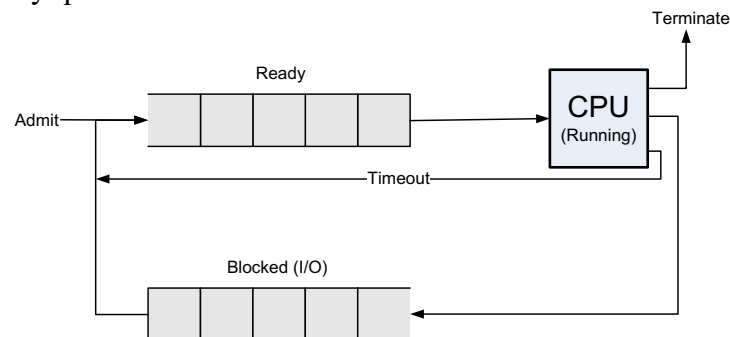
T3: 29 t.u. CPU burst, I/O request, 10 t.u. CPU burst, I/O request, 40 t.u. CPU burst, I/O request

To simplify the problem, assume that the time for running OS and Threading Library management code as negligible (i.e. time to switch threads or switch processes shall be consider to be 0).

- a) **Many to one model:** First consider that multi-threaded processes are supported by the computer system using a many to one model, that is, a thread library is responsible for managing threads in a process. Assume that the thread library uses two states (Runnable and Active) and a single scheduling queue as shown below for managing threads. A thread is moved to the Runnable state when it is waiting to become active and the Active state when it is allowed to execute (essentially assigned to the kernel thread). Typically other states are defined for managing user threads such as Sleeping and Stopped. But for simplicity, these states are not included for this problem. Assume that the Thread Library allows a threads maximum of 5 time units in the Active state before switching to another thread.



The OS on the other hand uses three states and two scheduling queues shown below. The Ready, Blocked, and Running states are used to represent processes ready to run, blocked waiting for an I/O operation to complete, and allocated to the CPU respectively. Note that states such as New, Terminated, and states for suspended processes are not included, again to keep the problem simple. Assume that the OS assigns the CPU to a process for a maximum of 10 time units, after which the CPU is allocated to the next process in the Ready queue.



Assume that both the threading library and the OS assign the Active/Running state to the thread/process at the head of the Runnable/Ready queues. Any process/thread added to these queues is placed at the end of the queue. This is a first-in first-out scheduling algorithm (this algorithm along with other scheduling algorithms shall be studied shortly in Module 4).

Complete the following table to show how processes and threads move between queues and the Active/Running state up to time 200 t.u. If a scheduling event at both user and kernel level occur at the same time, then assume that only the OS scheduling occurs, and that the user thread library will complete its scheduling operation the next time the CPU is assigned to the process (see time 50 for an example). Recall that the thread library allows a thread to be in the Active state for a maximum of 5 time units while processes are allocated to the CPU (i.e. in the running state) for a maximum of 10 time units. Also assume that time progresses for threads only when the process B is in the Running state, that is, allocated to the CPU. This means that when a thread is in the Active state and the process is in either the Ready or Blocked state, the threading library does not consider that time increases for the thread in the Active state.

[illegible]

* - the number in parentheses beside the A, B, T1, T2 and T3 represents the relative execution time (the time the process/thread executes on the CPU). Note that these times are incremented only when the process/thread is allocated to the CPU.

- b) **One to one model:** Now consider that the OS is aware of the threads and applies the three states not only to processes but also to threads in a multi-threaded process. No threading library is involved in the scheduling and management of the process threads. Assume that the OS allocates the CPU to a process/thread a maximum of 10 time units before switching the CPU to another processes/thread. Complete the following table to show how threads in process B and process A move between the Ready and Blocked queues and the Running state. As before assume a first-in, first-out queuing discipline for the Ready and Blocked queues. Like with case (a), assume that T2 arrives after T1 has executed 19 time units and T3 arrives after T1 has executed for 27 time units (i.e. it is T1 that invokes the threads T2 and T3).

[illegible]

c) On the following diagram, show which processes or threads are allocated to the CPU for both cases of (a) and (b). Note that the perspective presented is from the point of view of the OS – that is, in the case of the many to one model, indicate which processes are allocated to the CPU while in the case of the one to one model, show what threads of process B are allocated to the CPU. Comment on the difference in concurrency between the two models.

