

# Operating Systems Notes Chapter 7

July 28, 2025

## Contents

<b>1</b>	<b>Classic problems of synchronization</b>	<b>2</b>
1.1	Introduction . . . . .	2
1.2	Chapter objectives . . . . .	2
1.3	Classic problems of synchronization . . . . .	2
1.4	The bounded-buffer problem . . . . .	2
1.5	The readers-writers problem . . . . .	3
1.6	The dining-philosophers problem . . . . .	4
1.6.1	Semaphore solution . . . . .	4
1.6.2	Monitor solution . . . . .	4
1.7	Section glossary . . . . .	5
<b>2</b>	<b>Synchronization within the kernel</b>	<b>6</b>
2.1	Synchronization in Windows . . . . .	6
2.2	Synchronization in Linux . . . . .	6
2.3	Section glossary . . . . .	7
<b>3</b>	<b>POSIX synchronization</b>	<b>8</b>
3.1	POSIX mutex locks . . . . .	8
3.2	POSIX semaphores . . . . .	8
3.2.1	POSIX named semaphores . . . . .	8
3.2.2	POSIX unnamed semaphores . . . . .	9
3.3	POSIX condition variables . . . . .	9
3.4	Section glossary . . . . .	10
<b>4</b>	<b>Synchronization in Java</b>	<b>11</b>
4.1	Java monitors . . . . .	11
4.2	Block synchronization . . . . .	11
4.3	Reentrant locks . . . . .	13
4.4	Semaphores . . . . .	13
4.5	Condition variables . . . . .	13
4.6	Section glossary . . . . .	15
<b>5</b>	<b>Alternative approaches</b>	<b>16</b>
5.1	Transactional memory . . . . .	16
5.2	OpenMP . . . . .	17
5.3	Functional programming languages . . . . .	17
5.4	Section glossary . . . . .	18
<b>6</b>	<b>Summary</b>	<b>19</b>

# 1 Classic problems of synchronization

## 1.1 Introduction

- Previous chapter: Synchronization Tools.
- Focused on critical-section problem and race conditions with shared data.
- Examined tools to prevent race conditions:
  - Low-level hardware: memory barriers, compare-and-swap.
  - Higher-level: mutex locks, semaphores, monitors.
- Discussed challenges: liveness hazards like deadlocks.
- This chapter:
  - Applies synchronization tools to classic problems.
  - Explores synchronization mechanisms in Linux, UNIX, Windows.
  - Describes API details for Java and POSIX systems.

## 1.2 Chapter objectives

- Explain: bounded-buffer, readers-writers, dining-philosophers synchronization problems.
- Describe: specific tools used by Linux and Windows for process synchronization.
- Illustrate: how POSIX and Java solve process synchronization problems.
- Design and develop: solutions using POSIX and Java APIs.

## 1.3 Classic problems of synchronization

- Examples of concurrency-control problems.
- Used for testing new synchronization schemes.
- Solutions traditionally use semaphores; mutex locks can be used for binary semaphores in actual implementations.

## 1.4 The bounded-buffer problem

- Introduced in a previous chapter.
  - Illustrates power of synchronization primitives.
  - General structure presented; related programming project in exercises.
  - Producer and consumer processes share data structures:
- ```
int N;
semaphore mutex = 1;
semaphore empty = N;
semaphore full = 0
```
- Pool: N buffers, each holding one item.
  - `mutex`: binary semaphore, mutual exclusion for buffer pool access, initialized to 1.
  - `empty`: counts empty buffers, initialized to N.
  - `full`: counts full buffers, initialized to 0.
  - The producer process structure is as follows:

```
while (true) {
    . . .
    /* produce an item in next_produced */
    . . .
    wait(empty);
    wait(mutex);

    . . .
    /* add next_produced to the buffer */
    . . .
    signal(mutex);
    signal(full);
}
```

- The consumer process structure is as follows:

```
while (true) {
    wait(full);
    wait(mutex);

    . . .
    /* remove an item from buffer to next_consumed */
    . . .
    signal(mutex);
    signal(empty);

    . . .
    /* consume the item in next_consumed */
    . . .
}
```

```
}
```

- Symmetry between producer and consumer.
- Interpretation: producer produces full buffers for consumer, or consumer produces empty buffers for producer.

## 1.5 The readers-writers problem

- Shared database accessed by concurrent processes.
- **Readers**: only read database.
- **Writers**: update (read and write) database.
- Problem: If writer and another process (reader/writer) access simultaneously, chaos may ensue.
- Requirement: Writers must have exclusive access while writing.
- *Readers-writers problem*: Synchronization problem to ensure this.
- Variations involve priorities:
  - **First** readers-writers problem: No reader waits unless a writer already has permission. Readers don't wait for other readers if a writer is waiting.
  - **Second** readers-writers problem: Writer performs write ASAP once ready. If writer is waiting, no new readers may start.
- Starvation: Solutions may lead to starvation (writers in first case, readers in second).
- Solution to first readers-writers problem:
  - Shared data structures for reader processes:

```
semaphore rw_mutex = 1;
semaphore mutex = 1;
int read_count = 0;
```
  - `mutex` and `rw_mutex`: binary semaphores, initialized to 1.
  - `read_count`: integer, number of active readers, initialized to 0.
  - `rw_mutex`: common to reader and writer processes, acts as mutual exclusion for writers, used by first/last reader entering/exiting critical section.
  - `mutex`: ensures mutual exclusion when `read_count` is updated.
  - `read_count`: tracks current readers.
  - The writer process structure is as follows:

```
while (true) {
    wait(rw_mutex);

    . . .
    /* writing is performed */
    . . .
    signal(rw_mutex);
}
```
  - The reader process structure is as follows:

```
while (true) {
    wait(mutex);
    read_count++;
    if (read_count == 1)
        wait(rw_mutex);
    signal(mutex);

    . . .
    /* reading is performed */
    . . .
    wait(mutex);
    read_count--;
    if (read_count == 0)
        signal(rw_mutex);
    signal(mutex);
}
```
  - If writer in critical section and  $n$  readers waiting: 1 reader queued on `rw_mutex`,  $n - 1$  readers queued on `mutex`.
  - When writer executes `signal(rw_mutex)`, scheduler selects waiting readers or a single waiting writer.
- *Reader-writer locks*: generalization of problem/solutions.
  - Acquire lock by specifying mode: *read* or *write*.
  - Read mode: multiple processes concurrently.
  - Write mode: only one process (exclusive access).
  - Most useful when:
    - \* Easy to identify read-only vs. read-write processes.

- \* More readers than writers (increased concurrency compensates for overhead).

## 1.6 The dining-philosophers problem

- Five philosophers, circular table, five chairs, bowl of rice, five single chopsticks.
- Philosopher thinks, then gets hungry.
- Tries to pick up two closest chopsticks (left and right neighbors).
- Picks up one chopstick at a time. Cannot pick up if neighbor holds it.
- Eats with both chopsticks, then puts them down and thinks.
- Classic synchronization problem: example of allocating several resources among several processes.
- Goal: deadlock-free and starvation-free allocation.

### 1.6.1 Semaphore solution

- Each chopstick represented by a semaphore.
- Grab chopstick: `wait()` operation on semaphore.
- Release chopstick: `signal()` operation on semaphore.
- Shared data:
 

```
semaphore chopstick[5];
```
- All `chopstick` elements initialized to 1.
- The structure of philosopher  $i$  is as follows:
 

```
while (true) {
    wait(chopstick[i]);
    wait(chopstick[(i+1) % 5]);
    . . .
    /* eat for a while */
    . . .
    signal(chopstick[i]);
    signal(chopstick[(i+1) % 5]);
    . . .
    /* think for awhile */
    . . .
}
```
- Guarantees no two neighbors eat simultaneously.
- **Problem:** Could create deadlock.
  - \* Example: All five philosophers hungry, each grabs left chopstick. All `chopstick` elements become 0.
  - \* Each tries to grab right chopstick, delayed forever.
- Remedies to deadlock:
  - \* Allow at most four philosophers at table simultaneously.
  - \* Philosopher picks up both chopsticks only if both available (in critical section).
  - \* Asymmetric solution: odd-numbered philosopher picks left then right; even-numbered picks right then left.
- A previous chapter presents a deadlock-free solution.
- Satisfactory solution must guard against starvation (deadlock-free  $\neq$  starvation-free).

### 1.6.2 Monitor solution

- Deadlock-free solution using monitors.
- Restriction: Philosopher picks up chopsticks only if both available.
- Three states for a philosopher:
 

```
enum {THINKING, HUNGRY, EATING} state[5];
```
- `state[i] = EATING` only if neighbors not eating: `(state[(i+4) % 5] != EATING) and (state[(i+1) % 5] != EATING)`.
- Condition variable:
 

```
condition self[5];
```
- Allows philosopher  $i$  to delay if hungry but cannot get chopsticks.
- The distribution of chopsticks is controlled by the `DiningPhilosophers` monitor, which is defined below.
- Each philosopher  $i$  must invoke the operations `pickup()` and `putdown()` in the following sequence:
 

```
DiningPhilosophers.pickup(i);
    . . .
    eat
    . . .
    DiningPhilosophers.putdown(i);
```
- The monitor is defined as follows:

```

monitor DiningPhilosophers
{
    enum {THINKING, HUNGRY, EATING} state[5];
    condition self[5];

    void pickup(int i) {
        state[i] = HUNGRY;
        test(i);
        if (state[i] != EATING)
            self[i].wait();
    }

    void putdown(int i) {
        state[i] = THINKING;
        test((i + 4) % 5);
        test((i + 1) % 5);
    }

    void test(int i) {
        if ((state[(i + 4) % 5] != EATING) &&
            (state[i] == HUNGRY) &&
            (state[(i + 1) % 5] != EATING)) {
            state[i] = EATING;
            self[i].signal();
        }
    }

    initialization_code() {
        for (int i = 0; i < 5; i++)
            state[i] = THINKING;
    }
}

```

- Ensures no two neighbors eat simultaneously and no deadlocks.
- Possible for a philosopher to starve (solution not presented, left as exercise).

## 1.7 Section glossary

| Term                        | Definition                                                                                                                 |
|-----------------------------|----------------------------------------------------------------------------------------------------------------------------|
| readers-writers problem     | Synchronization problem where processes/threads either read or read/write shared data.                                     |
| reader-writer lock          | Lock for item access by read-only and read-write accessors.                                                                |
| dining-philosophers problem | Classic synchronization problem where multiple operators (philosophers) access multiple items (chopsticks) simultaneously. |

## 2 Synchronization within the kernel

### 2.1 Synchronization in Windows

- Windows OS: multithreaded kernel, supports real-time applications and multiple processors.
- Single-processor systems:
  - Kernel accesses global resource: temporarily masks interrupts for all interrupt handlers that may access the resource.
- Multiprocessor systems:
  - Kernel protects global resource access using spinlocks.
  - Spinlocks used only for short code segments.
  - Kernel ensures thread never preempted while holding a spinlock (for efficiency).
- Thread synchronization outside kernel: *dispatcher objects*.
  - Threads synchronize using: mutex locks, semaphores, events, timers.
  - Mutex locks: protect shared data; thread gains ownership to access, releases when finished.
  - Semaphores: behave as described in a previous chapter.
  - *Events*: similar to condition variables; notify waiting thread when condition occurs.
  - Timers: notify one (or more) threads when specified time expires.
- Dispatcher objects may be in one of two states:
  - A *signaled state*, which indicates that the object is available and a thread acquiring it will not block.
  - A *nonsignaled state*, which indicates that the object is not available and a thread attempting to acquire it will block.
- A mutex lock is acquired by a thread when it is in the signaled state, and it transitions to the nonsignaled state. When the thread releases the lock, it returns to the signaled state.
- Relationship between dispatcher object state and thread state:
  - Thread blocks on nonsignaled object: state changes from ready to waiting, placed in waiting queue.
  - Object moves to signaled state: kernel checks waiting threads.
  - Kernel moves one (or more) threads from waiting to ready state.
  - Number of threads selected depends on dispatcher object type:
    - \* Mutex: only one thread (mutex owned by single thread).
    - \* Event: all waiting threads.
- Mutex lock illustration:
  - Thread tries to acquire nonsignaled mutex: suspended, placed in waiting queue.
  - Mutex moves to signaled state (released by another thread): thread at front of queue moves from waiting to ready, acquires mutex.
- *Critical-section object*:
  - User-mode mutex, often acquired/released without kernel intervention.
  - Multiprocessor system: first uses spinlock while waiting.
  - If spins too long: acquiring thread allocates kernel mutex and yields CPU.
  - Efficient: kernel mutex allocated only when contention exists (rare in practice, significant savings).
- Programming project at end of chapter uses mutex locks and semaphores in Windows API.

### 2.2 Synchronization in Linux

- Prior to Version 2.6: nonpreemptive kernel (process in kernel mode could not be preempted).
- Now: Linux kernel is fully preemptive (task can be preempted while running in kernel).
- Mechanisms for synchronization in kernel:
  - **Atomic integers**:
    - \* Simplest synchronization technique.
    - \* Opaque data type: `atomic_t`.
    - \* All math operations are atomic (performed without interruption).
    - \* Example:

```
atomic_t counter;
int value;
```
  - \* Atomic operations:
    - `atomic_set(&counter,5);`: `counter = 5`
    - `atomic_add(10,&counter);`: `counter = counter + 10`
    - `atomic_sub(4,&counter);`: `counter = counter - 4`

- `atomic_inc(&counter);` `counter = counter + 1`
- `value = atomic_read(&counter);` `value = 12` (example result)
- \* Efficient for updating integer variables (e.g., counters); no locking overhead.
- \* Limited use: only for single integer variables. More sophisticated tools needed for multiple variables in race conditions.
- **Mutex locks:**
  - \* Protect critical sections within kernel.
  - \* Task invokes `mutex_lock()` before critical section, `mutex_unlock()` after.
  - \* If unavailable: task calling `mutex_lock()` sleeps, awakened when owner invokes `mutex_unlock()`.
- **Spinlocks and Semaphores:**
  - \* Linux also provides these (and reader-writer versions).
  - \* SMP machines: spinlock is fundamental locking mechanism, held for short durations.
  - \* Single-processor machines (e.g., embedded systems): spinlocks inappropriate. Replaced by enabling/disabling kernel preemption.
  - \* Summary for single-processor:
    - Instead of holding spinlock: kernel disables kernel preemption.
    - Instead of releasing spinlock: kernel enables kernel preemption.
- **Nonrecursive locks:**
  - Both spinlocks and mutex locks in Linux kernel are nonrecursive.
  - If thread acquires lock, cannot acquire same lock again without releasing it first.
  - Second attempt to acquire will block.
- **Disabling/Enabling kernel preemption:**
  - Linux approach: `preempt_disable()` and `preempt_enable()` system calls.
  - Kernel not preemptible if task running in kernel holds a lock.
  - Enforcement: each task has `thread-info` structure with `preempt_count` counter.
  - `preempt_count`: indicates number of locks held by task.
  - Lock acquired: `preempt_count` incremented.
  - Lock released: `preempt_count` decremented.
  - If `preempt_count > 0`: not safe to preempt kernel (task holds lock).
  - If `preempt_count = 0`: kernel can be safely interrupted (assuming no outstanding `preempt_disable()` calls).
- **When to use which lock:**
  - Spinlocks (and kernel preemption disable/enable): only when lock held for short duration.
  - Semaphores or mutex locks: when lock must be held for longer period.

## 2.3 Section glossary

| Term                    | Definition                                                                                                                                  |
|-------------------------|---------------------------------------------------------------------------------------------------------------------------------------------|
| dispatcher objects      | Windows scheduler feature controlling dispatching and synchronization. Threads synchronize via mutex locks, semaphores, events, and timers. |
| event                   | Windows OS scheduling feature, similar to a condition variable.                                                                             |
| critical-section object | User-mode mutex object in Windows OS, often acquired/released without kernel intervention.                                                  |

## 3 POSIX synchronization

- Synchronization methods in previous section: kernel-level, for kernel developers.
- POSIX API: available for user-level programmers, not part of specific OS kernel.
- Implemented using host OS tools.
- This section covers: mutex locks, semaphores, condition variables in Pthreads and POSIX APIs.
- Widely used for thread creation and synchronization on UNIX, Linux, macOS.

### 3.1 POSIX mutex locks

- Fundamental synchronization technique with Pthreads.
- Purpose: protect critical sections of code.
- Thread acquires lock before entering, releases upon exiting.
- Data type: `pthread_mutex_t`.
- Creation: `pthread_mutex_init()` function.
  - First parameter: pointer to mutex.
  - Second parameter: NULL for default attributes.
- Example:

```
#include <pthread.h>

pthread_mutex_t mutex;

/* create and initialize the mutex lock */
pthread_mutex_init(&mutex, NULL);
```
- Acquisition and Release: `pthread_mutex_lock()` and `pthread_mutex_unlock()`.
  - If `pthread_mutex_lock()` invoked and mutex unavailable: calling thread blocks until owner invokes `pthread_mutex_unlock()`.
- Protecting critical section example:

```
/* acquire the mutex lock */
pthread_mutex_lock(&mutex);

/* critical section */

/* release the mutex lock */
pthread_mutex_unlock(&mutex);
```
- Return values: 0 for correct operation, nonzero for error.

### 3.2 POSIX semaphores

- Many systems implementing Pthreads provide semaphores.
- Not part of POSIX standard; belong to POSIX SEM extension.
- Two types: *named* and *unnamed*.
- Differences: how they are created and shared between processes.
- Linux systems (Version 2.6+ kernel) support both.

#### 3.2.1 POSIX named semaphores

- Creation and opening: `sem_open()` function.
- Example:

```
#include <semaphore.h>
sem_t *sem;

/* Create the semaphore and initialize it to 1 */
sem = sem_open("SEM", O_CREAT, 0666, 1);
```
- "SEM": semaphore name.
- `O_CREAT` flag: semaphore created if it doesn't exist.
- `0666`: read and write access for other processes.
- Initialized to 1.
- Advantage: multiple unrelated processes can easily use common semaphore by name.
- Subsequent `sem_open()` calls (with same parameters) by other processes return descriptor to existing semaphore.
- Operations:
  - \* `wait()` → `sem_wait()`



```
* signal() → sem_post()
```

- Protecting critical section example:

```
/* acquire the semaphore */
sem_wait(sem);

/* critical section */

/* release the semaphore */
sem_post(sem);
```

- Supported by Linux and macOS.

### 3.2.2 POSIX unnamed semaphores

- Creation and initialization: `sem_init()` function.

- Parameters:

1. Pointer to the semaphore.
2. Flag indicating level of sharing.
3. Semaphore's initial value.

- Example:

```
#include <semaphore.h>
sem_t sem;

/* Create the semaphore and initialize it to 1 */
sem_init(&sem, 0, 1);
```

- Flag 0: semaphore shared only by threads in creating process.
- Nonzero flag: allows sharing between separate processes (by placing in shared memory).
- Initialized to 1.
- Operations: uses same `sem_wait()` and `sem_post()` as named semaphores.
- Protecting critical section example:

```
/* acquire the semaphore */
sem_wait(&sem);

/* critical section */

/* release the semaphore */
sem_post(&sem);
```

- Return values: 0 for success, nonzero for error.

## 3.3 POSIX condition variables

- Pthreads condition variables are similar to those described in a previous chapter.
- The difference is that the aforementioned chapter uses monitors for locking, whereas the C language, which is used with Pthreads, does not provide monitors.
- Locking accomplished by associating condition variable with a mutex lock.
- Data type: `pthread_cond_t`.
- Initialization: `pthread_cond_init()` function.
- Example of creating and initializing condition variable and associated mutex:

```
pthread_mutex_t mutex;
pthread_cond_t cond_var;

pthread_mutex_init(&mutex, NULL);
pthread_cond_init(&cond_var, NULL);
```

- Waiting on condition variable: `pthread_cond_wait()` function.
- Example of waiting for `a == b`:

```
pthread_mutex_lock(&mutex);
while (a != b)
    pthread_cond_wait(&cond_var, &mutex);
pthread_mutex_unlock(&mutex);
```

- Mutex lock must be acquired before `pthread_cond_wait()` call.
- Mutex protects data in conditional clause from race conditions.
- If condition not true: `pthread_cond_wait()` invoked, passing mutex and condition variable.

- `pthread_cond_wait()` releases mutex lock, allowing other threads to access/update shared data.
- Conditional clause within a loop: important to recheck condition after being signaled (protects against program errors).
- Signaling a condition variable: `pthread_cond_signal()` function.
- Thread modifying shared data invokes `pthread_cond_signal()` to signal one waiting thread.
- Example of signaling:
 

```
pthread_mutex_lock(&mutex);
a = b;
pthread_cond_signal(&cond_var);
pthread_mutex_unlock(&mutex);
```
- `pthread_cond_signal()` does NOT release mutex lock.
- Subsequent `pthread_mutex_unlock()` releases mutex.
- Once mutex released, signaled thread becomes owner of mutex and returns from `pthread_cond_wait()`.
- Programming problems/projects at end of chapter use Pthreads mutex locks, condition variables, and POSIX semaphores.

### 3.4 Section glossary

| Term              | Definition                                                                           |
|-------------------|--------------------------------------------------------------------------------------|
| named semaphore   | POSIX scheduling construct, exists in file system, shareable by unrelated processes. |
| unnamed semaphore | POSIX scheduling construct, usable only by threads in the same process.              |

## 4 Synchronization in Java

- Java language and API: rich support for thread synchronization since its origins.
- This section covers:
  - Java monitors (original mechanism).
  - Reentrant locks, semaphores, condition variables (introduced in Release 1.5).
- Focus on common locking and synchronization mechanisms.
- Java API has more features not covered (e.g., atomic variables, CAS instruction).

### 4.1 Java monitors

- Java provides a monitor-like concurrency mechanism, which is illustrated below with a `BoundedBuffer` class that implements the bounded-buffer problem using a monitor.
- The producer and consumer processes invoke the `insert()` and `remove()` methods, respectively.
- The structure of the `BoundedBuffer` class is as follows:

```
public class BoundedBuffer<E>
{
    private static final int BUFFER_SIZE = 5;

    private int count, in, out;
    private E[] buffer;

    public BoundedBuffer() {
        count = 0;
        in = 0;
        out = 0;
        buffer = (E[]) new Object[BUFFER_SIZE];
    }

    /* Producers call this method */
    public synchronized void insert(E item) {
        // details to be shown later
    }

    /* Consumers call this method */
    public synchronized E remove() {
        // details to be shown later
    }
}
```

- Every Java object has a single associated lock.
- **synchronized** method: entering requires owning the object's lock.
- Declared by placing **synchronized** keyword in method definition (e.g., `insert()`, `remove()`).
- Entering **synchronized** method:
  - Requires owning lock on `BoundedBuffer` object instance.
  - If lock owned by another thread: calling thread blocks, placed in object's *entry set*.
  - *Entry set*: set of threads waiting for lock to become available.
  - If lock available: calling thread becomes owner, enters method.
  - Lock released when thread exits method.
  - If entry set not empty on lock release: JVM arbitrarily selects thread from set to own lock (often FIFO in practice).
- The operation of the entry set is as follows: when a thread calls a **synchronized** method, it is added to the entry set for the object's lock. The thread is suspended until the lock is released, at which point the JVM selects a thread from the entry set to be granted the lock.
- Every object also has a *wait set* (initially empty).
- When thread enters **synchronized** method (owns lock):
  - May be unable to continue if condition not met (e.g., producer calls `insert()` and buffer is full).
  - Thread releases lock and waits until condition is met.

### 4.2 Block synchronization

- *Scope* of lock: time between acquisition and release.
- **synchronized** method: large scope if only small part manipulates shared data.
- Better: synchronize only the block of code manipulating shared data (smaller lock scope).
- Java allows block synchronization:

```

public void someMethod() {
    /* non-critical section */

    synchronized(this) {
        /* critical section */
    }

    /* remainder section */
}

```

- Only critical-section access requires ownership of **this** object lock.
- When thread calls **wait()** method:
  1. Releases lock for the object.
  2. Thread state set to blocked.
  3. Thread placed in wait set for the object.
- For example, when a producer thread invokes the **insert()** method and finds the buffer full, it calls **wait()**. This action releases the lock, blocks the producer, and places it in the wait set. A consumer thread can then acquire the lock, enter the **remove()** method, and free up space in the buffer.
- The relationship between the entry set and the wait set is as follows: when a thread in the wait set is notified, it is moved to the entry set and becomes eligible to be granted the lock.
- The full implementation of the **insert()** and **remove()** methods, which use **wait()** and **notify()**, is shown below:

```

/* Producers call this method */
public synchronized void insert(E item) {
    while (count == BUFFER_SIZE) {
        try {
            wait();
        }
        catch (InterruptedException ie) { }
    }

    buffer[in] = item;
    in = (in + 1) % BUFFER_SIZE;
    count++;

    notify();
}

/* Consumers call this method */
public synchronized E remove() {
    E item;

    while (count == 0) {
        try {
            wait();
        }
        catch (InterruptedException ie) { }
    }

    item = buffer[out];
    out = (out + 1) % BUFFER_SIZE;
    count--;

    notify();

    return item;
}

```

- **notify()** method:
  - Picks arbitrary thread **T** from wait set.
  - Moves **T** from wait set to entry set.
  - Sets state of **T** from blocked to runnable.
- **T** now eligible to compete for lock.
- Once **T** regains lock, returns from **wait()**, rechecks **count**.
- **notify()** ignored if no thread in wait set.
- Sequence of events with **wait()** and **notify()**:
  - Buffer full, lock available.
  - Producer calls **insert()**, enters, finds buffer full, calls **wait()**.
  - **wait()** releases lock, blocks producer, puts producer in wait set.
  - Consumer calls **remove()**, enters (lock available), removes item, calls **notify()**. Consumer still owns lock.
  - **notify()** moves producer from wait set to entry set, sets state to runnable.

- Consumer exits `remove()`, releases lock.
- Producer reacquires lock, resumes from `wait()`.
- Producer tests `while` loop, finds room, proceeds with `insert()`.
- Producer exits `insert()`, releases lock.
- `synchronized`, `wait()`, `notify()` are original Java mechanisms.
- Later Java API revisions introduced more flexible/robust locking.

### 4.3 Reentrant locks

- Simplest locking mechanism in API: `ReentrantLock`.
- Similar to `synchronized` statement: owned by single thread, provides mutual exclusive access to shared resource.
- Additional features: e.g., *fairness* parameter (favors longest-waiting thread).
- Acquisition: invoke `lock()` method.
  - If lock available OR invoking thread already owns it (*reentrant*): `lock()` assigns ownership, returns control.
  - If lock unavailable: invoking thread blocks until owner invokes `unlock()`.

- Implements `Lock` interface.

- Usage example:

```
Lock key = new ReentrantLock();
```

```
key.lock();
try {
    /* critical section */
}
finally {
    key.unlock();
}
```

- `try` and `finally` idiom:

- Ensures lock is released (via `unlock()`) after critical section completes or if exception occurs in `try` block.
- `lock()` not placed in `try` clause because it doesn't throw checked exceptions.
- Avoids `IllegalMonitorStateException` if unchecked exception occurs during `lock()` invocation (e.g., `OutOfMemoryError`), which would obscure original failure reason.

- `ReentrantLock` provides mutual exclusion.

- `ReentrantReadWriteLock`: Java API also provides this for scenarios with more readers than writers.

- Allows multiple concurrent readers but only one writer.

### 4.4 Semaphores

- The Java API provides a counting semaphore, as described in a previous chapter.
- Constructor: `Semaphore(int value)`.
- `value`: initial value of semaphore (negative allowed).
- `acquire()` method: throws `InterruptedException` if acquiring thread interrupted.
- Example using semaphore for mutual exclusion:

```
Semaphore sem = new Semaphore(1);
```

```
try {
    sem.acquire();
    /* critical section */
}
catch (InterruptedException ie) { }
finally {
    sem.release();
}
```

- `release()` placed in `finally` clause to ensure semaphore is released.

### 4.5 Condition variables

- Java API utility: condition variable.
- Functionality similar to `wait()` and `notify()` methods.
- Must be associated with a reentrant lock for mutual exclusion.
- Creation:
  1. Create a `ReentrantLock`.
  2. Invoke its `newCondition()` method.

- Returns a `Condition` object (representing condition variable for associated `ReentrantLock`).
- Example:
 

```
Lock key = new ReentrantLock();
Condition condVar = key.newCondition();
```
- Operations: `await()` and `signal()` methods.
- The function of these methods is the same as that of the `wait()` and `signal()` methods described in a previous chapter.
- Named vs. unnamed condition variables:
  - The monitors described in a previous chapter apply the `wait()` and `signal()` methods to *named* condition variables.
  - Java (language level): does not provide named condition variables.
  - Each Java monitor: associated with one unnamed condition variable.
  - `wait()` and `notify()` (Section Java monitors): apply only to this single unnamed condition variable.
  - When Java thread awakened via `notify()`: receives no info on why; reactivated thread must check condition itself.
  - Condition variables (this section): remedy this by allowing specific thread to be notified.
- Example: Five threads (0-4), shared variable `turn`.
- The `doWork(int threadNumber)` method demonstrates this concept.
  - In this example, only the thread whose `threadNumber` matches the shared variable `turn` is allowed to proceed; all other threads must wait.
  - The implementation of the `doWork()` method is as follows:
 

```
/* threadNumber is the thread that wishes to do some work */
public void doWork(int threadNumber)
{
    lock.lock();

    try {
        /**
         * If it's not my turn, then wait
         * until I'm signaled.
         */
        if (threadNumber != turn)
            condVars[threadNumber].await();

        /**
         * Do some work for awhile ...
         */

        /**
         * Now signal to the next thread.
         */
        turn = (turn + 1) % 5;
        condVars[turn].signal();
    }
    catch (InterruptedException ie) { }
    finally {
        lock.unlock();
    }
}
```
- Requires `ReentrantLock` and five condition variables (`condVars`).
- Initialization:
 

```
Lock lock = new ReentrantLock();
Condition[] condVars = new Condition[5];

for (int i = 0; i < 5; i++)
    condVars[i] = lock.newCondition();
```
- When thread enters `doWork()`:
  - If `threadNumber != turn`: invokes `await()` on its associated condition variable.
  - Resumes only when signaled by another thread.
  - After work: signals condition variable for next thread's turn.
- `doWork()` does not need to be `synchronized`.
  - `ReentrantLock` provides mutual exclusion.
  - `await()` on condition variable releases associated `ReentrantLock`.
  - `signal()` only signals condition variable; lock released by `unlock()`.
- Programming problems/projects at end of chapter use Pthreads mutex locks, condition variables, and POSIX semaphores.

## 4.6 Section glossary

| Term      | Definition                                                                              |
|-----------|-----------------------------------------------------------------------------------------|
| entry set | In Java, the set of threads waiting to enter a monitor.                                 |
| wait set  | In Java, a set of threads, each waiting for a condition that will allow it to continue. |
| scope     | The time between when a lock is acquired and when it is released.                       |

## 5 Alternative approaches

- Emergence of multicore systems: increased pressure to develop concurrent applications.
- Concurrent applications: increased risk of race conditions and liveness hazards (e.g., deadlock).
- Traditionally: mutex locks, semaphores, monitors used to address these.
- Challenge: as processing cores increase, designing multithreaded applications free from race conditions and deadlock becomes harder.
- This section explores: features in programming languages and hardware supporting thread-safe concurrent applications.

### 5.1 Transactional memory

- Idea originated in database theory, now used for process synchronization.
- *Memory transaction*: sequence of memory read-write operations that are atomic.
- If all operations complete: transaction committed.
- Otherwise: operations aborted and rolled back.
- Benefits obtained through language features.
- Example: `update()` function modifying shared data.
  - Traditional approach (with mutex locks/semaphores):

```
void update ()
{
    acquire();

    /* modify shared data */

    release();
}
```
  - Problems with traditional locking: deadlock, poor scalability with increasing threads (high contention for lock ownership).
  - Alternative: new language features using transactional memory.
  - Construct `atomic{S}`: ensures operations in `S` execute as a transaction.
  - Rewritten `update()` function:

```
void update ()
{
    atomic {
        /* modify shared data */
    }
}
```
  - Advantages of transactional memory:
    - \* Transactional memory system (not developer) guarantees atomicity.
    - \* No locks involved → deadlock not possible.
    - \* System can identify concurrent execution of statements in atomic blocks (e.g., concurrent read access).
    - \* Programmer identifying these situations (e.g., for reader-writer locks) becomes difficult as thread count grows.
- Implementation:
  - *Software transactional memory (STM)*:
    - \* Implemented exclusively in software; no special hardware needed.
    - \* Works by inserting instrumentation code inside transaction blocks (by compiler).
    - \* Manages transactions by examining concurrency and low-level locking needs.
  - *Hardware transactional memory (HTM)*:
    - \* Uses hardware cache hierarchies and cache coherency protocols.
    - \* Manages and resolves conflicts for shared data in separate processors' caches.
    - \* Requires no special code instrumentation (less overhead than STM).
    - \* Requires modification of existing cache hierarchies and cache coherency protocols.
- Status: existed for years without widespread implementation.
- Current trend: growth of multicore systems and emphasis on concurrent/parallel programming has prompted significant research.



## 5.2 OpenMP

- OpenMP supports parallel programming in a shared-memory environment.
- Includes: set of compiler directives and an API.
- `#pragma omp parallel`: compiler directive.
  - Code following this is a parallel region.
  - Performed by number of threads equal to processing cores.
- Advantage: OpenMP library handles thread creation and management (not application developers' responsibility).
- `#pragma omp critical`: compiler directive.
  - Specifies code region as a critical section.
  - Only one thread active at a time.
  - Ensures threads do not generate race conditions.
- Example: `update()` function modifying shared variable `counter`.

```
void update(int value)
{
    counter += value;
}
```

- If `update()` is part of/invoked from parallel region, race condition possible on `counter`.
- Remedy using critical-section compiler directive:

```
void update(int value)
{
    \#pragma omp critical
    {
        counter += value;
    }
}
```

- Behavior of critical-section directive:
  - Much like binary semaphore or mutex lock.
  - Ensures only one thread active in critical section at a time.
  - If thread tries to enter when another is active (owns section): calling thread blocks until owner exits.
  - Multiple critical sections: each can be named; rule specifies only one thread active in critical section of same name simultaneously.
- Advantages of OpenMP critical-section directive:
  - Generally considered easier to use than standard mutex locks.
- Disadvantages:
  - Developers must still identify possible race conditions.
  - Must adequately protect shared data using directive.
  - Deadlock still possible if two or more critical sections are identified (behaves like mutex lock).

## 5.3 Functional programming languages

- Most well-known languages (C, C++, Java, C#): *imperative* (or *procedural*) languages.
- Imperative languages:
  - Implement state-based algorithms.
  - Flow of algorithm crucial for correct operation.
  - State represented with variables and data structures.
  - Program state is mutable (variables can change values).
- Current emphasis on concurrent/parallel programming for multicore systems: greater focus on *functional* programming languages.
- Functional languages:
  - Different programming paradigm from imperative.
  - Fundamental difference: do not maintain state.
  - Once variable defined and assigned value, its value is immutable (cannot change).
  - Because mutable state disallowed: no concern with race conditions and deadlocks.
  - Most problems addressed in this chapter are nonexistent.
- Examples of functional languages:
  - Erlang: gained attention for concurrency support and ease of developing parallel applications.

- Scala: functional and object-oriented; syntax similar to Java and C#.

## 5.4 Section glossary

| Term                                       | Definition                                                                                                                                                      |
|--------------------------------------------|-----------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <b>transactional memory</b>                | Type of memory supporting memory transactions.                                                                                                                  |
| <b>memory transaction</b>                  | Sequence of memory read-write operations that are atomic.                                                                                                       |
| <b>software transactional memory (STM)</b> | Transactional memory implemented exclusively in software; no special hardware needed.                                                                           |
| <b>hardware transactional memory (HTM)</b> | Transactional memory using hardware cache hierarchies and cache coherency protocols to manage/resolve conflicts for shared data in separate processors' caches. |
| <b>imperative language</b>                 | Language for implementing state-based algorithms (e.g., C, C++, Java, C#).                                                                                      |
| <b>procedural language</b>                 | A language that implements state-based algorithms (e.g., C, C++, Java, C#).                                                                                     |
| <b>functional language</b>                 | Programming language that does not require states to be managed by programs written in it (e.g., Erlang, Scala).                                                |

## 6 Summary

- Classic process synchronization problems:
  - Bounded-buffer problem.
  - Readers-writers problem.
  - Dining-philosophers problem.
- Solutions use tools from "Synchronization Tools" chapter:
  - Mutex locks.
  - Semaphores.
  - Monitors.
  - Condition variables.
- **Windows synchronization:**
  - Uses dispatcher objects.
  - Uses events to implement synchronization tools.
- **Linux synchronization:**
  - Uses various approaches to protect against race conditions.
  - Includes atomic variables.
  - Includes spinlocks.
  - Includes mutex locks.
- **POSIX API synchronization:**
  - Provides mutex locks.
  - Provides semaphores.
  - Provides condition variables.
  - Two forms of semaphores:
    - \* Named semaphores: easily accessed by unrelated processes by name.
    - \* Unnamed semaphores: cannot be shared as easily; require placement in shared memory.
- **Java synchronization:**
  - Rich library and API for synchronization.
  - Available tools:
    - \* Monitors (provided at language level).
    - \* Reentrant locks (supported by API).
    - \* Semaphores (supported by API).
    - \* Condition variables (supported by API).
- **Alternative approaches to critical-section problem:**
  - Transactional memory.
  - OpenMP.
  - Functional languages.
- **Functional languages:**
  - Intriguing alternative programming paradigm to procedural languages.
  - Unlike procedural languages, do not maintain state.
  - Generally immune from race conditions and critical sections.