

# Operating Systems Notes - Chapter 13

July 27, 2025

## Contents

<b>1</b>	<b>File concept</b>	<b>2</b>
<b>2</b>	<b>Access methods</b>	<b>6</b>
<b>3</b>	<b>Directory structure</b>	<b>8</b>
3.1	Single-level directory . . . . .	8
3.2	Two-level directory . . . . .	8
3.3	Tree-structured directories . . . . .	8
3.4	Acyclic-graph directories . . . . .	9
3.5	General graph directory . . . . .	9
<b>4</b>	<b>Protection</b>	<b>10</b>
4.1	Types of access . . . . .	10
4.2	Access control . . . . .	10
4.3	Other protection approaches . . . . .	11
<b>5</b>	<b>Memory-mapped files</b>	<b>12</b>
5.1	Basic mechanism . . . . .	12
5.2	Shared memory in the Windows API . . . . .	12
<b>6</b>	<b>Summary</b>	<b>13</b>

# 1 File concept

## Things to learn

- Explain file system function.
- Describe file system interfaces.
- Discuss file system design tradeoffs: access methods, sharing, locking, directory structures.
- Explore file system protection.

## Introduction

- **File**: collection of related information defined by its creator.
- OS maps files onto physical mass-storage devices.
- File system: describes how files map to physical devices, how accessed/manipulated.
- File systems designed for efficient access (physical storage can be slow).
- Other requirements: file sharing support, remote access.
- File system: most visible OS aspect for users.
- Provides mechanism for on-line storage/access to OS data/programs and user data.
- Consists of:
  - Collection of files (storing related data).
  - Directory structure (organizes, provides info about files).
- Most file systems live on storage devices (nonvolatile, persistent).

## File Concept

- Computers store info on various media (NVM, HDDs, tapes, optical disks).
- OS provides uniform logical view of stored info.
- OS abstracts physical properties to define logical storage unit: the **file**.
- Files mapped by OS onto physical devices.
- Storage devices usually nonvolatile, so contents persistent.
- File: named collection of related information recorded on secondary storage.
- User perspective: smallest allotment of logical secondary storage (data written only within a file).
- Commonly: files represent programs (source/object) and data.
- Data files: numeric, alphabetic, alphanumeric, binary.
- Files can be free form (text) or rigidly formatted.
- Generally: file is sequence of bits, bytes, lines, or records; meaning defined by creator/user.
- File concept is extremely general.
- Use stretched beyond original confines (e.g., UNIX ‘proc’ file system uses file-system interfaces for system info access).
- Info in file defined by creator.
- File has defined structure, depends on type.
- **Text file**: sequence of characters organized into lines.
- **Source file**: sequence of functions (declarations, executable statements).
- **Executable file**: series of code sections loader can bring into memory and execute.

## File Attributes

- File named for human users, referred by name (e.g., ‘example.c’).
- Some systems differentiate case, others not.
- File independent of creator process, user, system.
- Example: ‘example.c’ created by one user, edited by another, copied to USB/email/network, still ‘example.c’.
- Second copy independent if no sharing/synchronization.
- File attributes vary by OS, typically include:
  - **Name**: symbolic, human-readable.
  - **Identifier**: unique tag (number), non-human-readable, identifies file within file system.
  - **Type**: for systems supporting different file types.
  - **Location**: pointer to device and file location on device.
  - **Size**: current size (bytes, words, blocks), possibly max allowed.

- **Protection:** access-control info (read, write, execute).
- **Timestamps and user identification:** creation, last modification, last use (useful for protection, security, monitoring).
- Newer file systems support **extended file attributes:** character encoding, file checksum.
- **File info window:** GUI view of file metadata (e.g., macOS).
- Info about all files kept in directory structure, on same device as files.
- Directory entry: file name, unique identifier (locates other attributes).
- Directory size can be large (MBs/GBs).
- Directories stored on device, brought into memory piecemeal.

## File Operations

- File is an abstract data type.
- OS provides system calls: create, write, read, reposition, delete, truncate files.
- **Creating a file:**
  1. Find space in file system.
  2. Make entry for new file in directory.
- **Opening a file:**
  - All operations except create/delete require ‘open()’ first.
  - Returns file handle used as argument in other calls.
- **Writing a file:**
  - System call: open file handle, info to write.
  - System searches directory for file location.
  - System keeps **write pointer** to next write location (sequential).
  - Write pointer updated after each write.
- **Reading a file:**
  - System call: file handle, memory location for next block.
  - Directory searched for entry.
  - System keeps **read pointer** to next read location (sequential).
  - Read pointer updated after each read.
  - Current operation location: per-process **current-file-position pointer** (shared by read/write).
- **Repositioning within a file:**
  - Current-file-position pointer repositioned to given value.
  - No actual I/O involved.
  - Also known as file **seek**.
- **Deleting a file:**
  - Search directory for named file.
  - Release all file space for reuse.
  - Erase/mark directory entry as free.
  - **Hard links:** multiple names for same file; actual content deleted only when last link deleted.
- **Truncating a file:**
  - Erase contents but keep attributes.
  - Reset file length to zero, release file space.
- These 7 are minimal set. Other common: appending, renaming.
- Primitive operations combine for others (e.g., copy file).
- Operations to get/set file attributes (e.g., length, owner).
- To avoid constant directory searching: ‘open()’ system call before first use.
- OS keeps **open-file table:** info about all open files.
- File specified by index into table (no searching).
- File closed: OS removes entry from open-file table, releases locks.
- ‘create()’ and ‘delete()’ work with closed files.
- Some systems implicitly open/close files (job termination).
- Most systems require explicit ‘open()’/‘close()’.
- ‘open()’: takes file name, searches directory, copies entry to open-file table.

- ‘open()’ accepts access-mode info (create, read-only, read-write, append-only).
- Mode checked against file permissions. If allowed, file opened.
- ‘open()’ returns pointer to open-file table entry; used in all I/O ops.
- ‘open()’/‘close()’ complicated with simultaneous opens by multiple processes.
- OS uses two levels of internal tables:
  - Per-process table: tracks files process has open, current file pointer, access rights, accounting.
  - System-wide open-file table: process-independent info (disk location, access dates, size).
- File opened by another process: new entry in process’s table points to system-wide entry.
- Open-file table has **open count**: number of processes with file open.
- ‘close()’ decreases count; when zero, entry removed.
- File locks: allow one process to lock file/sections, prevent others.
- Useful for shared files (e.g., system log).
- **Shared lock**: multiple processes acquire concurrently (like reader lock).
- **Exclusive lock**: only one process at a time (like writer lock).
- Not all OS provide both types.
- **Mandatory** vs. **advisory** file-locking mechanisms.
- Mandatory: OS prevents other processes from accessing locked file (e.g., Windows).
- Advisory: OS does not prevent access; text editor must manually acquire lock (e.g., UNIX).
- Mandatory: OS ensures locking integrity. Advisory: developers ensure locks.
- File locks require same precautions as process synchronization (e.g., hold exclusive locks only during access, avoid deadlock).

## File Types

- OS should recognize/support file types for reasonable operations.
- Example: prevent outputting binary-object program as garbage.
- Common technique: include type as part of file name (name.extension).
- User/OS can tell type from name (e.g., ‘resume.docx’, ‘server.c’).
- OS uses extension to indicate type and allowed operations (e.g., ‘.com’, ‘.exe’, ‘.sh’ for execution).
- ‘.sh’ is a **shell script** (ASCII commands).
- Application programs use extensions (e.g., Java compilers expect ‘.java’).
- Extensions not always required; user can omit, app looks for expected extension.
- Extensions are “hints” to applications, not OS-enforced.
- macOS: each file has type (‘.app’), creator attribute (program that created it).
- Creator attribute set by OS during ‘create()’ call, enforced.
- Example: word processor file has word processor as creator; double-click opens app, loads file.
- UNIX: **magic number** at beginning of some binary files indicates data type (e.g., image format).
- Text magic number for text files (e.g., shell language).
- Not all files have magic numbers; system features not solely based on this.
- UNIX does not record creating program.
- UNIX allows file-name-extension hints, but not enforced/depended on by OS; aid users.
- File types indicate internal structure.
- Source/object files: structures match expectations of programs reading them.
- Certain files must conform to OS-understood structure (e.g., executable file structure for loading/running).
- Some OS extend this: system-supported file structures with special operations.
- Disadvantage of OS supporting multiple file structures: large, cumbersome OS.
- OS needs code for each supported structure.
- Every file may need to be defined as one of supported types.
- New applications with unsupported structures: problems.
- Example: encrypted file not ASCII text, not executable binary. May need to circumvent/misuse OS file-type mechanism.
- Some OS impose minimal file structures (UNIX, Windows).
- UNIX: each file is sequence of 8-bit bytes; no OS interpretation.
- Provides maximum flexibility but little support; application must interpret structure.

- All OS must support at least one structure: executable file.

## Internal File Structure

- Internally, locating offset within file can be complicated for OS.
- Disk systems: well-defined block size (sector size).
- All disk I/O in units of one block (physical record); all blocks same size.
- Physical record size unlikely to match desired logical record length.
- Logical records may vary in length.
- Solution: packing logical records into physical blocks.
- UNIX: all files are streams of bytes. Each byte individually addressable by offset.
- File system automatically packs/unpacks bytes into physical disk blocks (e.g., 512 bytes/block).
- Logical record size, physical block size, packing technique determine records per block.
- Packing by user app or OS. File considered sequence of blocks.
- Basic I/O functions operate in terms of blocks.
- Conversion from logical records to physical blocks: simple software problem.
- Disk space allocated in blocks: some portion of last block wasted (**internal fragmentation**).
- Example: 512-byte blocks, 1,949-byte file → 4 blocks (2,048 bytes); 99 bytes wasted.
- All file systems suffer internal fragmentation; larger block size → greater fragmentation.

## Section glossary

Term	Definition
<b>file</b>	Smallest logical storage unit; collection of related information.
<b>text file</b>	File containing text (alphanumeric characters).
<b>source file</b>	File containing program source code.
<b>executable file</b>	File containing program ready for loading/execution.
<b>extended file attributes</b>	Extended metadata (character encoding, checksums).
<b>file info window</b>	GUI view of file metadata.
<b>write pointer</b>	Location in file for next write.
<b>read pointer</b>	Location in file for next read.
<b>current-file-position pointer</b>	Per-process pointer to next read/write location.
<b>seek</b>	Operation of changing current file-position pointer.
<b>hard links</b>	File-system links where file has two+ names to same inode.
<b>open-file table</b>	OS data structure with details of every open file.
<b>open count</b>	Number of processes with an open file.
<b>shared lock</b>	File lock allowing concurrent acquisition by multiple processes.
<b>exclusive lock</b>	File lock allowing only one process to acquire at a time.
<b>advisory file-lock mechanism</b>	File-locking system where OS does not enforce locking.
<b>shell script</b>	File containing set series of commands specific to shell.
<b>magic number</b>	Number at start of file indicating data type.
<b>internal fragmentation</b>	Wasted disk space in last block of file due to block allocation.

## 2 Access methods

### Accessing File Information

- Files store information; must be accessed and read into memory.
- Information accessed in several ways.
- Some systems provide only one access method; others many.
- Choosing right method: major design problem.

### Sequential Access

- Simplest access method: **sequential access**.
- Information processed in order, one record after another.
- Most common (editors, compilers).
- `read_next()`: reads next portion, automatically advances file pointer.
- `write_next()`: appends to end, advances to end of newly written material.
- File can be reset to beginning.
- Some systems: skip forward/backward `n` records.
- Based on tape model of file.
- Works on sequential-access devices and random-access ones.

### Direct Access

- Another method: **direct access** (or **relative access**).
- File: fixed-length **logical records**.
- Programs read/write records rapidly in no particular order.
- Based on disk model of file (disks allow random access).
- File viewed as numbered sequence of blocks/records.
- No restrictions on read/write order.
- Great use for immediate access to large info amounts (e.g., databases).
- Example: airline reservation system, flight info in block identified by flight number.
- Direct-access file operations: include block number as parameter.
- `read(n)`, `write(n)` instead of `read_next()`, `write_next()`.
- Alternative: retain `read_next()`, `write_next()`, add `position_file(n)`.
- Block number provided by user: **relative block number**.
- Relative block number: index relative to beginning of file (first is 0, next 1, etc.).
- OS decides file placement (**allocation problem**).
- Prevents user from accessing non-file portions of file system.
- Some systems start relative block numbers at 0, others at 1.
- Satisfying request for record `N`: turned into I/O request for `N` bytes starting at `N * (logical record length)`.
- Logical records fixed size: easy to read, write, delete a record.
- Not all OS support both sequential and direct access.
- Some require file defined as sequential/direct at creation.
- Simulate sequential access on direct-access file: keep `cp` variable for current position.
- Simulating direct-access on sequential-access: extremely inefficient and clumsy.

### Other Access Methods

- Built on top of direct-access method.
- Involve constructing an **index** for the file.
- Index: contains pointers to various blocks (like book index).
- Find record: search index, use pointer to access file directly.
- Example: retail-price file (UPCs, prices). Sorted by UPC.
- Index: first UPC in each block. Can be kept in memory.
- Binary search index → find block → access block.
- Large files: index file too large for memory.
- Solution: index for the index file (primary index → secondary index → data).

- Example: IBM ISAM (indexed sequential-access method).
- Small main index points to disk blocks of secondary index.
- Secondary index blocks point to actual file blocks.
- File sorted on key.
- Find item: binary search main index → get secondary index block → binary search secondary index → find block with record → sequential search block.
- Any record located by at most two direct-access reads.

Section glossary

Term	Definition
sequential access	File-access method: contents read in order, beginning to end.
direct access	File-access method: contents read in random order.
relative access	File-access method: contents read in random order.
logical records	File contents logically designated as fixed-length structured data.
relative block number	Index relative to beginning of file (first is block 0).
allocation problem	OS determination of where to store file blocks.
index	Access method built on direct access; file contains index with pointers to contents.

## 3 Directory structure

Directory: symbol table translating file names to file control blocks. Organization must allow:

- Insert entries
- Delete entries
- Search for named entry
- List all entries

Operations on a directory:

- **Search for a file:** Find entry for particular file; find files matching pattern.
- **Create a file:** Add new files to directory.
- **Delete a file:** Remove file from directory; may leave hole, defragmentation needed.
- **List a directory:** List files and their entry contents.
- **Rename a file:** Change file name when contents/use changes; may change position.
- **Traverse the file system:** Access every directory/file; for backup or space release.

### 3.1 Single-level directory

- Simplest structure: all files in same directory.
- Easy to support and understand.
- Limitations:
  - Files must have unique names (name collision problem for multiple users).
  - Difficult for single user to remember many file names.

### 3.2 Two-level directory

- Separate directory for each user.
- Each user has own **user file directory (UFD)**.
- System's **main file directory (MFD)** indexed by user name/account, points to UFD.
- When user refers to file, only their UFD searched.
- Different users can have same file names (unique within each UFD).
- OS searches only user's UFD for create/delete.
- UFDs created/deleted by special system program (restricted to administrators).
- Disadvantages:
  - Isolates users; disadvantage for cooperation.
  - To access another user's file, must specify user name and file name.
  - Two-level directory as a tree (MFD root, UFDs descendants, files leaves).
  - User name + file name = **path name**.
  - Example: `/userb/test.txt` or `C:\userb\test`.
- System files:
  - Copying system files to each UFD wastes space.
  - Solution: special user directory for system files (e.g., user 0).
  - OS first searches local UFD, then special system directory.
  - Sequence of directories searched: **search path**.
  - Search path can be extended; users can have own search paths.

### 3.3 Tree-structured directories

- Generalization of two-level directory to arbitrary height.
- Most common directory structure.
- Root directory; every file has unique path name.
- Directory (or subdirectory) contains files or subdirectories.
- Directory often treated as special file; one bit defines entry as file (0) or subdirectory (1).
- Special system calls to create/delete directories.
- Each process has a **current directory**.
- Reference to file: current directory searched.
- If not in current directory: specify path name or change current directory.



- Initial current directory from accounting file.
- Path names:
  - **Absolute path name:** begins at root (e.g., `"/`), follows path down.
  - **Relative path name:** defines path from current directory.
  - Example: if current is `/spell/mail`, `prt/first` is same as `/spell/mail/prt/first`.
- User defines subdirectories for organization (e.g., by topic, info type).
- Deletion of a directory:
  - If empty: entry simply deleted.
  - If not empty:
    - \* Some systems: only delete if empty (user must delete contents recursively first).
    - \* Others (e.g., UNIX `rm -r`): delete directory and all its files/subdirectories recursively. More convenient, but dangerous.
- Users can access other users' files by specifying path name or changing current directory.

### 3.4 Acyclic-graph directories

- Allows directories to share subdirectories and files.
- No cycles (loops) in the graph.
- Shared file: one actual file exists, changes visible to all.
- Shared subdirectory: new files appear in all shared subdirectories.
- Implementation:
  - **Link:** pointer to another file/subdirectory (e.g., absolute/relative path name).
  - **Resolve:** use path name in link to locate real file.
  - OS ignores links during directory traversal to preserve acyclic structure.
  - Duplicate all info in both sharing directories: entries identical, but consistency issues on modification.
- Problems:
  - Multiple absolute path names for same file (aliasing).
  - Traversing entire file system: avoid traversing shared structures more than once.
  - Deletion: when can space be deallocated?
    - \* Deleting file leaves dangling pointers if other links exist.
    - \* Symbolic links: deletion of link doesn't affect original file. If original deleted, links dangle.
    - \* Preserve file until all references deleted: use **reference count**.
    - \* Increment count on new link/entry, decrement on deletion. Delete when count is 0.
    - \* UNIX uses for **hard links**.

### 3.5 General graph directory

- Allows cycles in the directory structure.
- Primary advantage of acyclic graph: simpler traversal and deletion algorithms.
- Problems with cycles:
  - Infinite loops during search/traversal.
  - Reference count may not be 0 even if file/directory is inaccessible.
  - Requires **garbage collection** to determine when space can be reallocated (time consuming for disk-based systems).
- Avoiding cycles: computationally expensive to detect.
- Simpler: bypass links during directory traversal.

Table 1: Section glossary

Term	Definition
<b>user file directory (UFD)</b>	Per-user directory of files in two-level directory implementation.
<b>main file directory (MFD)</b>	Index pointing to each UFD in two-level directory implementation.
<b>path name</b>	File-system name for a file, containing mount-point and directory-entry info to locate it (e.g., "C:/foo/bar.txt").
<b>search path</b>	Sequence of directories searched for an executable file when a command is executed.
<b>absolute path name</b>	Path name starting at the top of the file system hierarchy.
<b>relative path name</b>	Path name starting at a relative location (e.g., current directory).
<b>acyclic graph</b>	Directory structure implementation that contains no cycles (loops).
<b>link</b>	File that has no contents but points to another file.
<b>resolve</b>	To follow a link and find the target file.
<b>hard links</b>	File-system links where a file has two or more names pointing to the same inode.
<b>garbage collection</b>	Recovery of space containing no-longer-valid data.

## 4 Protection

Information safety:

- Physical damage: **reliability** (duplicate copies, backups).
- Improper access: **protection**.

Protection mechanisms:

- User name/password authentication.
- Encrypting secondary storage.
- Firewalling network access.
- Multiuser systems: advanced mechanisms for valid data access.

### 4.1 Types of access

- Need for controlled access.
- Protection limits types of file access.
- Operations controlled:
  - **Read**: Read from file.
  - **Write**: Write or rewrite file.
  - **Execute**: Load and execute file.
  - **Append**: Write new info at end of file.
  - **Delete**: Delete file, free space.
  - **List**: List name and attributes.
  - **Attribute change**: Change file attributes.
- Higher-level functions (rename, copy, edit) often implemented by system programs using lower-level calls. Protection at lower level.

### 4.2 Access control

- Access dependent on user identity.
- Most general scheme: **access-control list (ACL)**.
- ACL specifies user names and allowed access types.
- OS checks ACL; allows if listed, denies otherwise.
- Advantages: complex access methodologies.
- Disadvantages:
  - Lengthy lists (tedious to construct, especially if users unknown).
  - Variable-size directory entries (complicated space management).
- Condensed ACL: three user classifications:
  - **Owner**: User who created file.
  - **Group**: Set of users sharing file, needing similar access.
  - **Other**: All other users.
- Common approach: combine ACLs with owner, group, universe scheme (e.g., Solaris).
- UNIX permissions:
  - Three fields: owner, group, universe.
  - Each field: three bits **rwX** (read, write, execute).

- **r** for read, **w** for write, **x** for execution.
- Example: **rw**x for owner, **rw-** for group, **r--** for others.
- **d** as first character indicates subdirectory.
- Sample listing shows links, owner, group, size, date, name.
- Combining ACLs and permissions:
  - User interface challenge: how to show optional ACLs.
  - Solaris: "+" appended to regular permissions (e.g., **-rw-r--r--+**).
  - Commands like **setfacl** and **getfacl** manage ACLs.
  - Windows: GUI for ACL management.
  - Precedence: ACLs typically take precedence over group permissions (specificity priority).

### 4.3 Other protection approaches

- Password with each file:
  - Effective if passwords random and changed often.
  - Disadvantages: many passwords to remember; single password for all files (all-or-none protection).
  - Some systems: password with subdirectory.
  - More commonly: encryption of partition/files, with key password management.
- Directory protection in multilevel structures:
  - Control creation/deletion of files in directory.
  - Control user's ability to determine file existence (listing directory contents).
  - If path name refers to file, user needs access to both directory and file.
  - Different access rights depending on path name in acyclic/general graphs.

Table 2: Section glossary

Term	Definition
access-control list	A list of user names allowed to access a file.

## 5 Memory-mapped files

Alternative file access method: **memory mapping** a file.

- Treat file I/O as routine memory accesses using virtual memory techniques.
- Can lead to significant performance increases.

### 5.1 Basic mechanism

- Map disk block to page(s) in memory.
- Initial access: demand paging, page fault.
- Page-sized portion of file read into physical page.
- Subsequent reads/writes: handled as routine memory accesses.
- Simplifies and speeds up file access by avoiding `read()` and `write()` system call overhead.
- Writes to memory-mapped file not necessarily immediate to secondary storage.
- Generally, updates written back when file closed.
- Under memory pressure, intermediate changes may go to swap space.
- Some OS (e.g., Solaris) memory-map all file I/O, even with standard calls, to kernel address space.
- Multiple processes can map same file concurrently for data sharing.
- Writes by one process visible to others mapping same section.
- Implemented by virtual memory map pointing to same physical page.
- Supports copy-on-write: processes share read-only, get own copies for modification.
- Processes use mutual exclusion for shared data coordination.
- Shared memory often implemented by memory mapping files.

### 5.2 Shared memory in the Windows API

- Outline for shared memory using memory-mapped files:
  1. Create a **file mapping** for the file.
  2. Establish a **view** of the mapped file in process's virtual address space.
- Second process opens and creates view of same mapped file.
- Mapped file acts as shared-memory object for inter-process communication.
- Example: Producer writes, Consumer reads.
- Steps:
  1. Open file with `CreateFile()` (returns `HANDLE`).
  2. Create file mapping with `CreateFileMapping()` (uses file `HANDLE`).
  3. Establish view with `MapViewOfFile()` (uses mapped object `HANDLE`).
- `CreateFileMapping()` creates a **named shared-memory object** (e.g., `SharedObject`).
- `MapViewOfFile()` returns pointer to shared-memory object; accesses to this memory are accesses to the file.
- Entire file or portion can be mapped.
- Mapped file may be demand-paged.
- Both processes remove view with `UnmapViewOfFile()`.

Table 3: Section glossary

Term	Definition
<b>memory mapping</b>	File-access method where file is mapped into process memory space for direct memory access
<b>file mapping</b>	In Windows, the first step in memory-mapping a file.
<b>view</b>	In Windows, an address range mapped in shared memory; second step in memory-mapping a file.
<b>named shared-memory object</b>	In Windows API, a section of a memory-mapped file accessible by name from multiple processes.

## 6 Summary

- File: abstract data type, sequence of logical records (byte, line, complex data). OS may support record types or leave to application.
- OS task: map logical file concept to physical storage (hard disk, NVM). May order logical records into physical records.
- Directories: organize files.
  - Single-level directory: naming problems in multiuser systems (unique names required).
  - Two-level directory: separate directory for each user, solves naming problems. Lists file name, location, length, type, owner, times.
  - Tree-structured directory: generalization of two-level, allows subdirectories for organization.
  - Acyclic-graph directory: allows sharing of subdirectories/files, but complicates searching/deletion.
  - General graph structure: complete flexibility in sharing, but may require garbage collection for unused space.
- Remote file systems: challenges in reliability, performance, security. Distributed information systems manage user, host, access info for shared state.
- File protection: needed on multiuser systems.
  - Access controlled by type: read, write, execute, append, delete, list directory.
  - Protection via access lists, passwords, other techniques.